

Today

- Likelihood inference concepts
- Likelihood training algorithm

p-values

- Use constraint; prefer uncertainty intervals
- Some key points (of many)
 - p-value is not the probability that: “null is true”, “data were generated by the null” or “by chance alone”
 - $p < 0.05$ does not mean “the null hypothesis is false”
 - small p-value does not mean “the effect was large or important”
 - $p > 0.05$ does not mean “there was no effect”, or “the null is true”, or “the effect was small”
 - if many replicated studies have $p > 0.05$ it does not provide accumulated evidence of “no effect”

Likelihood in data science

- This week: **pure** likelihood inference
 - Learning goal: understand likelihood
- Likelihood is also used in
 - Frequentist: as a sample statistic
 - Bayesian: part of the posterior
 - Information theory: e.g. AIC
 - likelihood + complexity penalty

Likelihood

We put all the **assumptions** about the data generating process in the **likelihood function**

Likelihood principle

All the **evidence** in an observation (data) about the parameters (model) is in the **likelihood function**

Likelihood function

Counts all the ways the data could have happened **for a given model or hypothesis**

Reading


McElreath Ch 2

Figs 2.2 - 2.4

Marbles in a bag

We know: 4 marbles, 2 colors, marbles drawn randomly with replacement

Goal: what is in the bag?



hypothesis, model

H_1, M_1


H_2, M_2

H_3, M_3

H_4, M_4

H_5, M_5

Data






FIGURE 2.2. The 64 possible paths generated by assuming the bag contains one blue and three white marbles.

i.e. assuming we have H_2 , M_2

[●○○○]

$$L = 3/64$$




FIGURE 2.3. After eliminating paths inconsistent with the observed sequence, only 3 of the 64 paths remain.

Paths for data



given M_2






FIGURE 2.3. After eliminating paths inconsistent with the observed sequence, only 3 of the 64 paths remain.

What is the likelihood for
2 blue + 1 white *in any order*?

given M_2

[●○○○]

$$L = 9/64$$




FIGURE 2.3. After eliminating paths inconsistent with the observed sequence, only 3 of the 64 paths remain.

What is the likelihood for
2 blue + 1 white *in any order*?

given M_2

[●○○○]

The likelihood

- Probability of the data **given a model**

$$L = P(\text{Data} \mid \text{Model 2})$$

Data

Model 2

"given"

A conditional probability

The likelihood

- Probability of the data **given a model**

$$L = P(\text{Data} \mid \text{Model 2})$$

Data Model 2

"given"

A conditional probability

$$P(y \mid M_2) = P(y \mid \theta_2)$$

$$y = ["b", "w", "b"]$$

could be a vector

θ indicates parameters
(number of blue & white)

Likelihood of model or H

A model is more likely than another if it is the model for which the **data** are more probable

Notice that this doesn't mention the *probability* of the model, only the **probability of the data**

Inference: likelihood ratio

$$\frac{P(y|\theta_2)}{P(y|\theta_1)}$$


Strength of evidence for model 2
compared to model 1

Models

Pathways

Likelihood

M_1



?

M_2



3

3/64

M_3



?

M_4



?

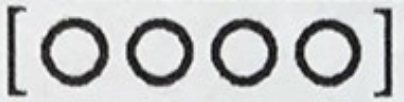




M_5




?

Data =



	Models	Pathways	Likelihood
M_1		0	0
M_2		3	$3/64$
M_3		?	
M_4		?	
M_5		0	0

Data =



M_2 [●○○○]

$L = 3/64$


M_4 [●●●○]

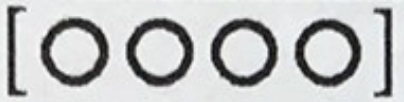




$L = 9/64$

M_3 [●●○○]


$L = 8/64$

Data =



	Models	Pathways	Likelihood
M_1		0	0
M_2		3	$3/64$
M_3		8	$8/64$
M_4		9	$9/64$
M_5		0	0

Data =



$$\frac{P(y|\theta_4)}{P(y|\theta_2)} = \frac{9}{3} = 3$$

$$\frac{P(y|\theta_4)}{P(y|\theta_3)} = \frac{9}{8} = 1.125$$

$$\frac{P(y|\theta_3)}{P(y|\theta_2)} = \frac{8}{3} = 2.\dot{6}$$

Notes

- Not frequentist
- Not the same or even similar to a sampling distribution
 - we have not invoked multiple repeated samples
 - probability of the data, not probability of a sample statistic

Likelihood inference for the linear model

Model algorithm

Writing down the model (DGP):

$$y_i \sim \text{Normal}(\mu_i, \sigma)$$

Stochastic process

$$\mu_i = \beta_0 + \beta_1 x_i$$

Deterministic process

Model algorithm

Writing down the model:

$$y_i \sim \text{Normal}(\mu_i, \sigma) \quad \text{or} \quad y \sim \text{Normal}(\mu, \sigma)$$
$$\mu_i = \beta_0 + \beta_1 x_i \quad \mu = \beta_0 + \beta_1 x$$

Model algorithm

Writing down the model:





$$y_i \sim \text{Normal}(\mu_i, \sigma) \quad \text{or} \quad y \sim \text{Normal}(\mu, \sigma)$$
$$\mu_i = \beta_0 + \beta_1 x_i \quad \mu = \beta_0 + \beta_1 x$$

Algorithm (vectorized):


```
lin_skel <- function(beta_0, beta_1, x) {  
  return(beta_0 + beta_1 * x)  
}  
y_stoch <- function(mu=lin_skel(beta_0, beta_1, x), sigma) {  
  return(rnorm(n=length(mu), mean=mu, sd=sigma))  
}
```


Simulating the model


```
y_stoch(mu=lin_skel(beta_0=300, beta_1=-9, x=x), sigma=30)
```



Likelihood (linear, Normal)



Likelihood (linear, Normal)



Likelihood (linear, Normal)

Likelihood for the model:

$$\begin{aligned} L(\theta) &= P(y|\theta) = P(y|\beta_0, \beta_1, \sigma) \\ &= P(y|\beta_0, \beta_1, \sigma, x) \end{aligned}$$

```
dnorm(y, mean=lin_skel(beta_0, beta_1, x), sd)
```


Total likelihood for a data set

One data point: $P(y_7|\theta)$

All data points: $\prod_i^n P(y_i|\theta)$

because probabilities multiply together to give total probability (n is the number of datapoints). **Independence** is assumed.

Likelihood (linear, Normal)



Support function

The log likelihood:

$$\sum_i^n \ln P(y_i|\theta)$$

Instead of multiplying small probabilities, it is more accurate and convenient to sum their logs.

```
sum(dnorm(y, mean=lin_skel(beta_0, beta_1, x), sd, log=TRUE))
```

Training algorithm: Maximum likelihood

The **values of the parameters** that **maximize the likelihood**. In other words, the model that maximizes the probability of the data.

An **optimization** problem.

In practice: minimize the negative log likelihood. The model with the most support, has the smallest negative log likelihood.

Training algorithm

```
lin_skel <- function(beta_0, beta_1, x) {  
  return(beta_0 + beta_1 * x)  
}
```

call the linear model

```
lm_nll <- function(p, y, x) {  
  mu <- lin_skel(beta_0=p[1], beta_1=p[2], x=x)  
  nll <- -sum(dnorm(y, mean=mu, sd=p[3], log=TRUE))  
  return(nll)  
}
```

minus sum log likelihoods

```
fit_lm <- optim(p=start_pars, lm_nll, y=y, x=x)  
fit_lm
```

Find maximum
likelihood estimates
(MLE) for all 3
parameters

Training general approach

- 1) process model function (deterministic skeleton)
 - the biology
 - > `biomod(parameters)`
- 2) nll function (error distribution)
 - the stochasticity or error
 - > `-sum(ddist(data, dpar=biomod, error_parameters, log=T))`
- 3) optimize
 - find biology parameters and error parameters
 - > `optim(parameters, nllfunc, data)`
- This recipe is the same no matter how complicated the process model or what the error distribution is