

Today

- Finish off last lecture
- Bayesian training algorithms

McElreath Ch8

- Learning goals:
- Understand and use MCMC algorithms to sample from the posterior distribution
- Recognize and fix bad sampling scenarios
- Use HMC implemented in Stan via R packages

Main points McElreath Ch8

- MCMC: Monte Carlo Markov Chain
- Series of random numbers where each number depends on the previous one
- Sample less from low probability areas; more bang for your random buck
- Algorithms
 - Metropolis-Hastings
 - Gibbs sampling
 - HMC: Hamiltonian Monte Carlo
- We'll mostly use HMC

MCMC algorithms

Algorithm (general)

for many iterations

propose new value for parameter

calculate the probability of accepting the proposal:

$$P_{\text{accept}} = \min(\Pr(\text{proposal}) / \Pr(\text{current}), 1)$$

accept proposal randomly with $\text{Bern}(P_{\text{accept}})$

plot posterior distribution (histogram) of parameter values

where $\Pr() = \text{prior} \times \text{likelihood}$

Rosenbluth algorithm

aka Metropolis-Hastings

Algorithm (original)

for many iterations

propose new value for parameter:

draw Unif(-max_d, max_d)

proposal = current parameter + draw

calculate the probability of accepting the proposal:

$P_{\text{accept}} = \min(\Pr(\text{proposal}) / \Pr(\text{current}), 1)$

accept proposal randomly with $\text{Bern}(P_{\text{accept}})$

plot posterior distribution (histogram) of parameter values

where $\Pr() = \text{prior} \times \text{likelihood}$

MCMC algorithms

- Two other important algorithms
- Gibbs sampling
 - needs **conjugate** priors (so less general)
 - prior such that the posterior is the same as the prior
 - e.g. norm prior \times norm lik = norm posterior
 - beta prior \times binom lik = beta posterior
 - Hamiltonion Monte Carlo (HMC)

MCMC algorithms

- Get an intuition for their behavior:
- <https://chi-feng.github.io/mcmc-demo/app.html#HamiltonianMC,standard>

Bayesian tools

- Stan
 - Gelman group
 - Hamiltonian Monte Carlo
 - Betancourt (2017) A conceptual introduction to Hamiltonian Monte Carlo (<https://arxiv.org/abs/1701.02434A>)
 - open source
 - models with continuous parameters only
 - state of the art
 - R and Python interfaces (packages)
- <http://mc-stan.org>

Bayesian tools

- BUGS (Bayesian inference Using Gibbs Sampling) (and Metropolis-Hastings)
- <http://www.openbugs.info>
- Older, original standard tool for MCMC
- Exceedingly difficult to run on Mac
- Many newer tools are **based on BUGS** code style
- Lots of books and publications use BUGS
- Recommend: need to know about historically but don't use anymore

Bayesian tools

- JAGS (Just Another Gibbs Sampler)
- <http://mcmc-jags.sourceforge.net/>
 - cross platform, open source
 - basically the same as BUGS
 - often faster
 - recommended for models that can't be fit in Stan (e.g. discrete parameters)
 - easy install
- Best to run from R
 - Install R2jags package (install from R)

Bayesian tools

- Others:
- Nimble
 - somewhat common in ecology
- Julia: Turing package
- Python: PyMC (rapid progression)

Main points McElreath Ch8

- Using HMC via Stan to fit models
- Now getting posterior samples from HMC
- Use ulam in rethinking to do HMC to follow examples
- Same syntax as sampost

Main points McElreath Ch8

ulam or sampost

```
m1 <- ulam(  
  alist(  
    y ~ dnorm(mu, sigma),  
    mu <- a + b * x,  
    a ~ dnorm(0, 100),  
    b ~ dnorm(0, 10),  
    sigma ~ dcauchy(0, 2)  
  ),  
  data=d1)
```

Main points McElreath Ch8

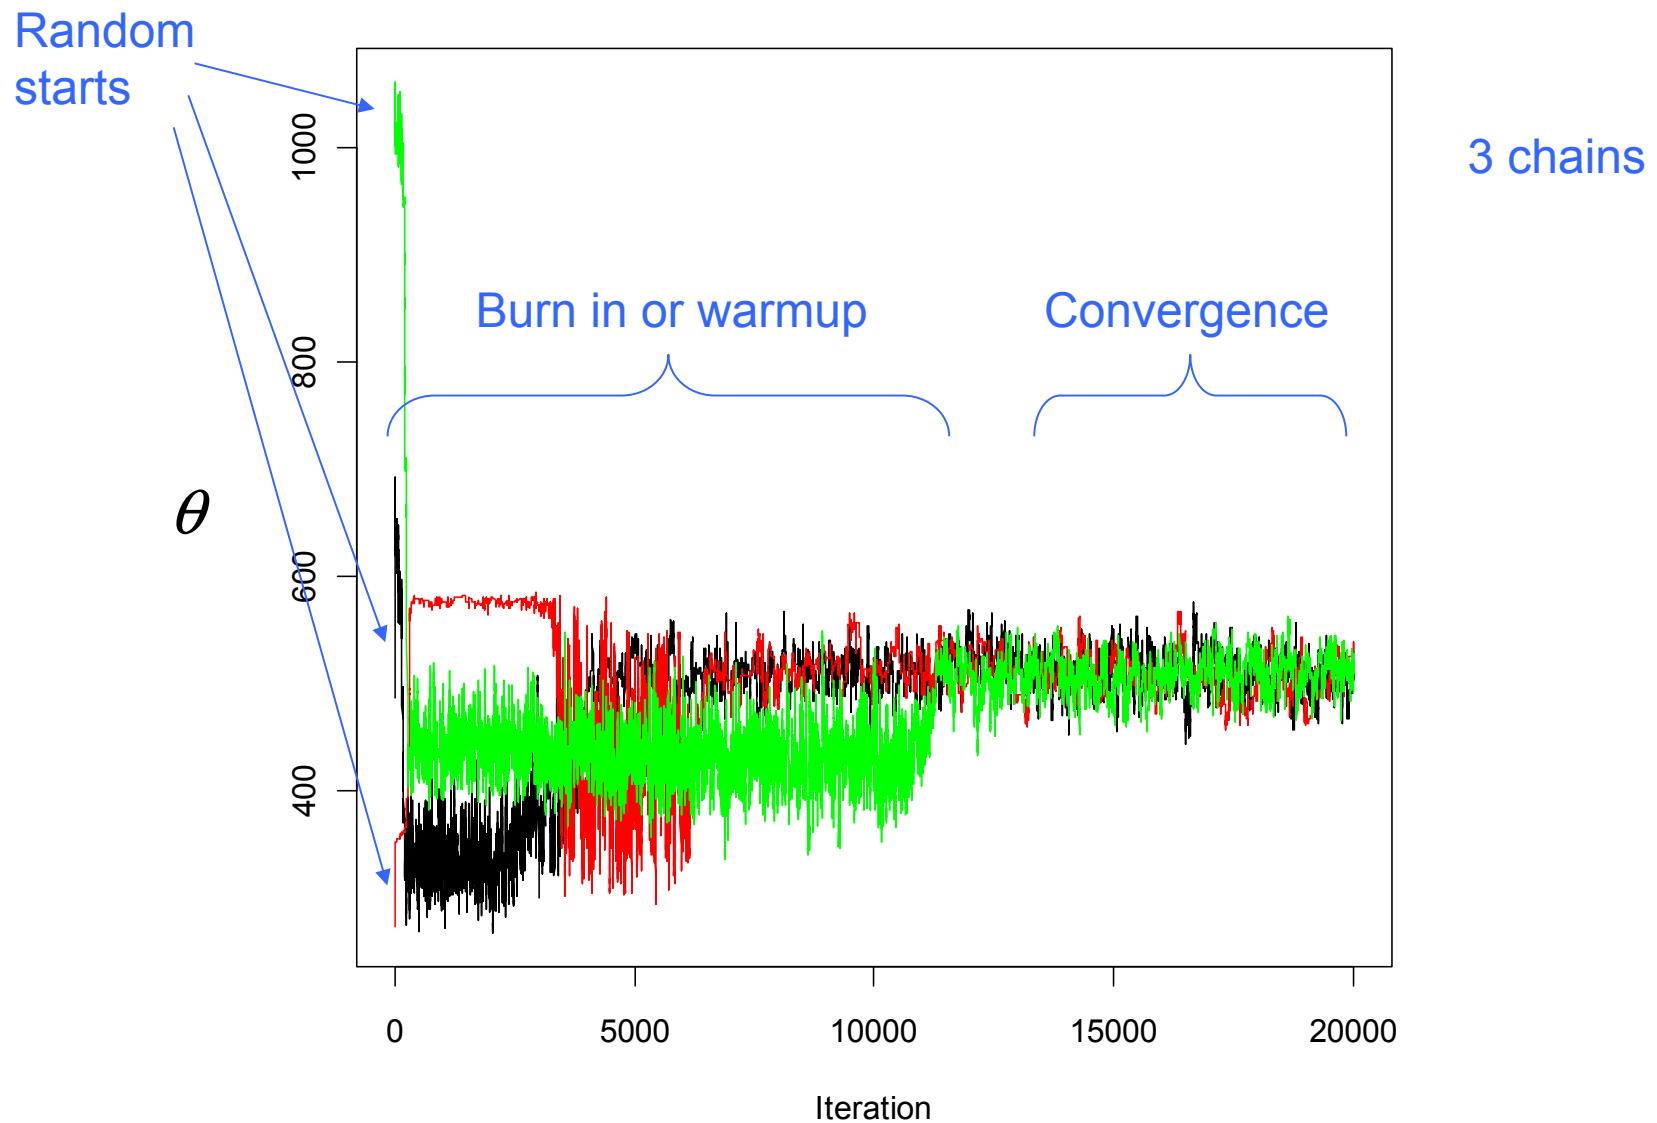
ulam or sampost

```
m1 <- ulam(  
  alist(  
    y ~ dnorm(mu, sigma),  
    mu <- a + b * x,  
    a ~ dnorm(0, 100),  
    b ~ dnorm(0, 10),  
    sigma ~ dcauchy(0, 2)  
  ),  
  data=d1)
```

Main points McElreath Ch8

- Good choice of priors (**weakly informative**) can be helpful to tame model fit
 - e.g. **Half-Cauchy** instead of uniform
- MCMC **diagnostics** to judge convergence of fit
 - `rhat`, `n_eff`
 - plot chain traces ("time series")
- **Visualize** posteriors
 - histograms, pairs plot

Chains



Bad

Good

Chains

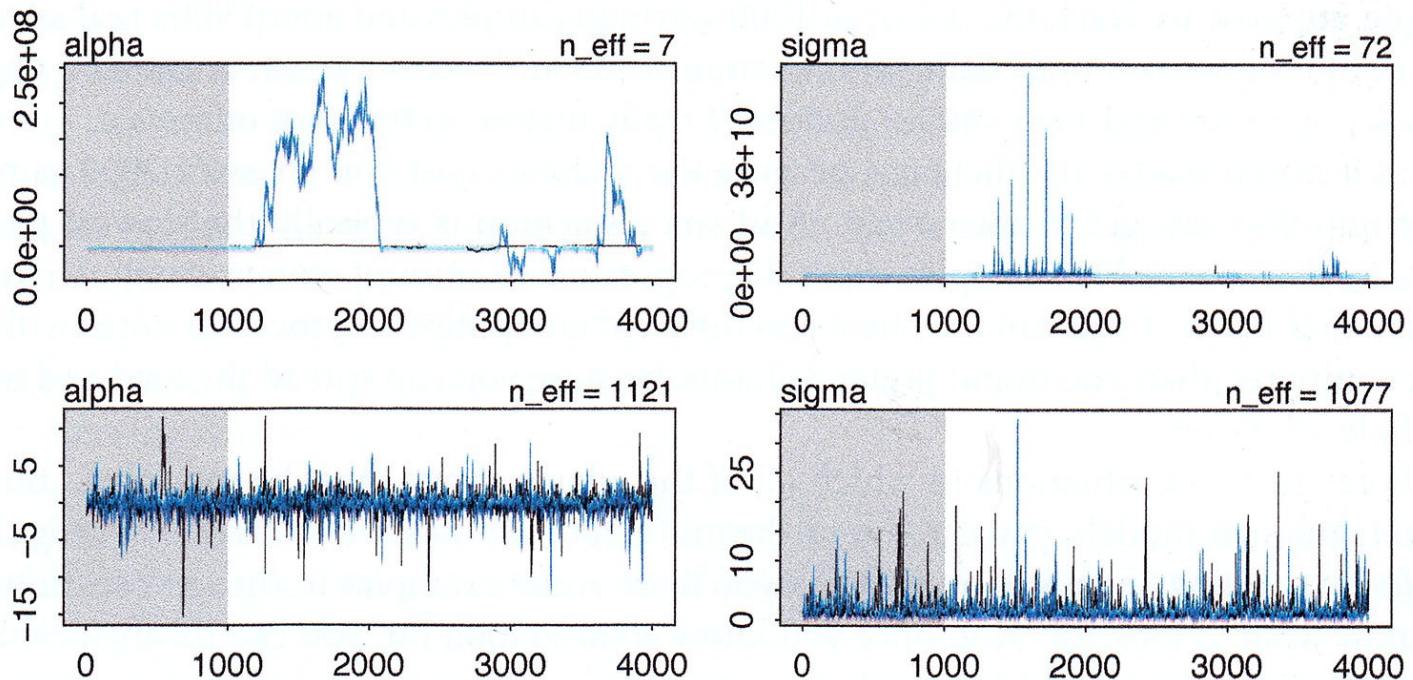


Fig 8.5

Chains

Not converged Converged

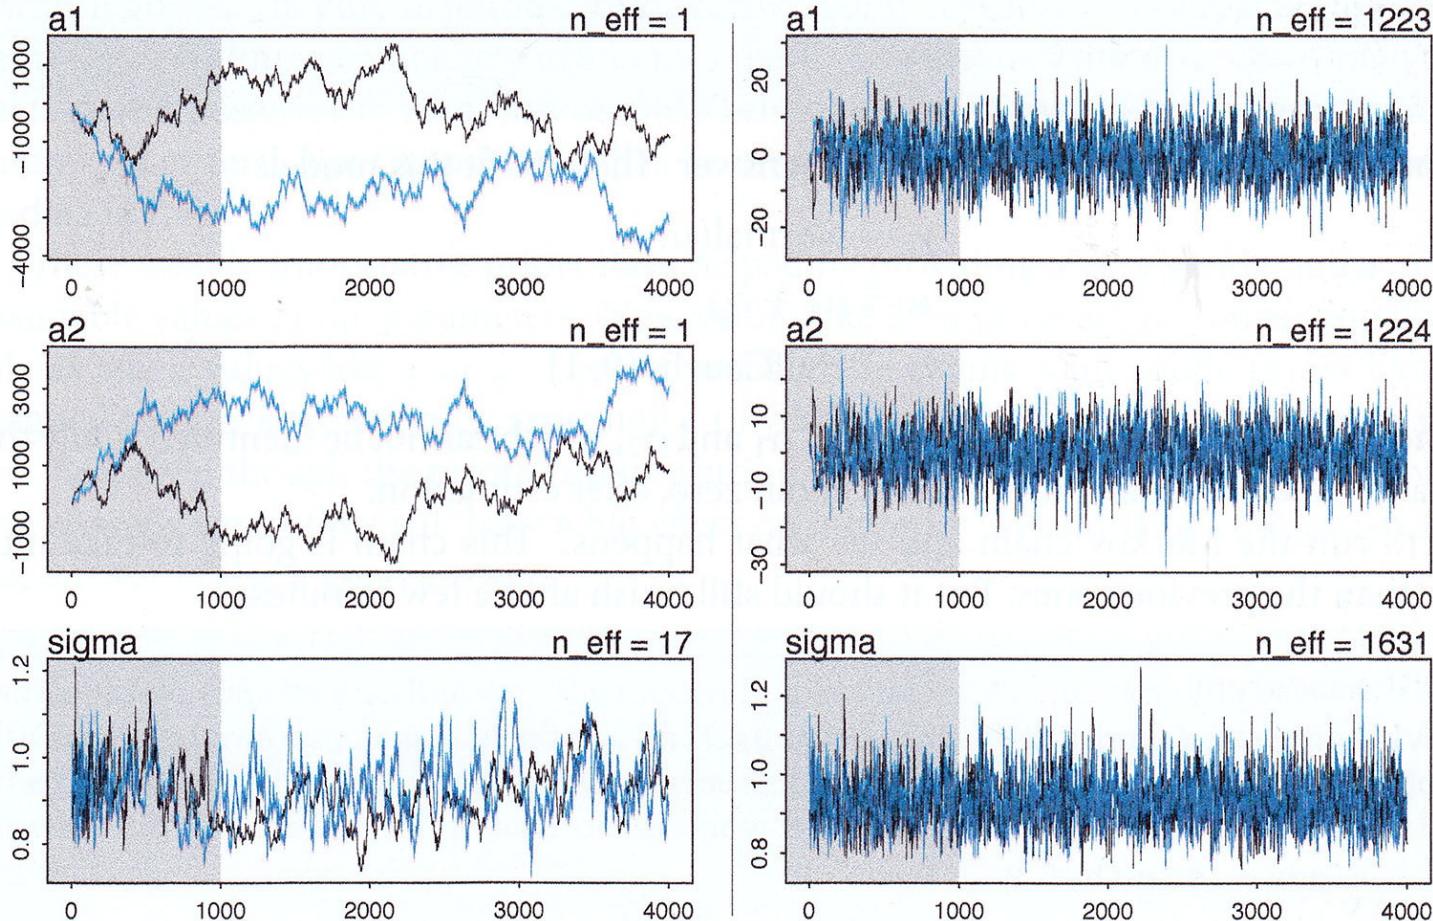


Fig 8.7

How to fix

- Better starting values
- Weakly informative priors
- Uncorrelated parameters (e.g. standardized)
- Less common: adjust MCMC algorithm parameters